

Fungi Edibility Predictive Image Classification Model - Final Notebook

PROGRAMMER: Alex Karadjov

-

Imports

```
In [1]: 1 import csv
        2 import os
        3 os.environ["PROTOCOL_BUFFERS_PYTHON_IMPLEMENTATION"] = "python"
        4 import yaml
        5 import joblib
        6
        7 import pandas as pd
        8 import numpy as np
        9 import pickle
       10 import streamlit as st
       11 import matplotlib.pyplot as plt
       12 import tensorflow as tf
       13 from keras import models
       14 from PIL import Image
       15
       16
       17 from tensorflow.keras import layers
       18 from tensorflow.keras.preprocessing.image import img_to_array, load_image
       19 from tensorflow.keras.utils import to_categorical
       20 from tensorflow.keras.models import load_model
       21 from tensorflow.keras.preprocessing import image
       22 from tensorflow.keras.applications.vgg16 import preprocess_input
       23
       24 from sklearn.pipeline import Pipeline
       25 from sklearn.preprocessing import FunctionTransformer
       26 from sklearn.model_selection import train_test_split
       27 from imblearn.over_sampling import SMOTE
       28 from sklearn.dummy import DummyClassifier
       29 from sklearn.metrics import accuracy_score
       30 from sklearn.preprocessing import LabelEncoder
```

-

Importing images and labels from previous notebook

csv file needs to be reshaped

```
In [ ]: 1 images_flat = np.loadtxt('images.csv', delimiter=',')
        2 images = images_flat.reshape((images.shape[0],) + images.shape[1:])
```

```
In [ ]: 1 labels_encoded = np.loadtxt('labels.csv', delimiter=',')
        2 labels = label_encoder.inverse_transform(labels_encoded)
```

```
In [ ]: 1 print(images.shape)
        2 print(labels_decoded.shape)
```

-

Training data for models.

```
In [15]: 1 train_images, test_images, train_labels, test_labels = train_test_split
```

```
In [16]: 1 (train_images, train_labels), (test_images, test_labels) = tf.keras.
```

-

Setting train and test values to correct formatting for modeling.

```
In [17]: 1 train_images = train_images.reshape((60000, 28, 28, 1))
        2 train_images = train_images.astype('float32') / 255.0
        3
        4 test_images = test_images.reshape((10000, 28, 28, 1))
        5 test_images = test_images.astype('float32') / 255.0
        6
        7
        8 train_images = tf.convert_to_tensor(train_images)
        9 train_labels = tf.convert_to_tensor(train_labels)
```

Testing shapes.

```
In [18]: 1 train_images.shape
```

```
Out[18]: TensorShape([60000, 28, 28, 1])
```

-

```
In [19]: 1 train_labels.shape
```

```
Out[19]: TensorShape([60000])
```

Final conversion so I can apply data set to different models.

```
In [20]: 1 train_labels = to_categorical(train_labels)
```

-

Trying Dataset with Dummy Classifier

Very simple first model to try out.

```
In [21]: 1 dummy_model = DummyClassifier(strategy="uniform")
```

```
In [22]: 1 dummy_model.fit(train_images, train_labels)
```

```
Out[22]: DummyClassifier(strategy='uniform')
```

```
In [23]: 1 predicted_labels = dummy_model.predict(test_images)
2
3
4 predicted_labels = np.argmax(predicted_labels, axis=1)
5
6
7 accuracy = accuracy_score(test_labels, predicted_labels)
8 print("Accuracy:", accuracy)
```

Accuracy: 0.1062

Interesting not sure this is correct or working...

-

Trying Neural Network Model

Will be used as baseline model final model will contain more layers.

```
In [24]: 1 baseline_model = tf.keras.Sequential([
2     layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 3)),
3     layers.SeparableConv2D(filters=64, kernel_size=(3, 3), activation='relu'),
4     layers.Flatten(),
5     layers.Dense(64, activation='relu'),
6     layers.Dense(10, activation='softmax')
7 ])
```

-

```
In [25]: 1 baseline_model.compile(optimizer='adam', loss='categorical_crossentropy')
2
3 baseline_model.fit(train_images, train_labels, epochs=10, batch_size=128)
```

```
Epoch 1/10
1875/1875 [=====] - 38s 20ms/step - loss: 0.1408 - accuracy: 0.9574
Epoch 2/10
1875/1875 [=====] - 34s 18ms/step - loss: 0.0458 - accuracy: 0.9856
Epoch 3/10
1875/1875 [=====] - 38s 20ms/step - loss: 0.0289 - accuracy: 0.9903
Epoch 4/10
1875/1875 [=====] - 40s 21ms/step - loss: 0.0194 - accuracy: 0.9936
Epoch 5/10
1875/1875 [=====] - 40s 21ms/step - loss: 0.0145 - accuracy: 0.9953
Epoch 6/10
1875/1875 [=====] - 39s 21ms/step - loss: 0.0110 - accuracy: 0.9964
Epoch 7/10
1875/1875 [=====] - 42s 23ms/step - loss: 0.0091 - accuracy: 0.9971
Epoch 8/10
1875/1875 [=====] - 40s 21ms/step - loss: 0.0073 - accuracy: 0.9974
Epoch 9/10
1875/1875 [=====] - 42s 22ms/step - loss: 0.0074 - accuracy: 0.9976
Epoch 10/10
1875/1875 [=====] - 42s 23ms/step - loss: 0.0044 - accuracy: 0.9986
```

```
Out[25]: <tensorflow.python.keras.callbacks.History at 0x194fff865e0>
```

Really good baseline model, I probably wont have to do much more in final model.

-

Final Model Training and Testing

For my final model I added a few more layers to improve accuracy

```
In [26]: ▶ 1 model = tf.keras.Sequential([
2     layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
3     layers.SeparableConv2D(filters=64, kernel_size=(3, 3), activation='relu'),
4     layers.MaxPooling2D((2, 2)),
5     layers.Conv2D(filters=128, kernel_size=(3, 3), activation='relu'),
6     layers.GlobalAveragePooling2D(),
7     layers.Flatten(),
8     layers.Dense(64, activation='relu'),
9     layers.Dense(10, activation='softmax')
10 ])
11
12 early_stopping = tf.keras.callbacks.EarlyStopping(
13     monitor="val_loss",
14     patience=4,
15     verbose=1,
16     restore_best_weights=True
17 )
```

```
In [ ]: ▶ 1
```

```
In [ ]: ▶ 1
```

```
In [27]: ▶ 1 model.compile(optimizer='adam', loss='categorical_crossentropy', metr  
2  
3 model.fit(train_images, train_labels, epochs=10, batch_size=32, callt
```

Epoch 1/10
1874/1875 [=====>.] - ETA: 0s - loss: 0.4832 - accuracy: 0.8428WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 54s 29ms/step - loss: 0.4830 - accuracy: 0.8428

Epoch 2/10
1875/1875 [=====] - ETA: 0s - loss: 0.1407 - accuracy: 0.9574 ETA: 0s - loss: 0.1411 - WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 54s 29ms/step - loss: 0.1407 - accuracy: 0.9574

Epoch 3/10
1875/1875 [=====] - ETA: 0s - loss: 0.0901 - accuracy: 0.9722WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 56s 30ms/step - loss: 0.0901 - accuracy: 0.9722

Epoch 4/10
1875/1875 [=====] - ETA: 0s - loss: 0.0702 - accuracy: 0.9788WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 57s 30ms/step - loss: 0.0702 - accuracy: 0.9788

Epoch 5/10
1875/1875 [=====] - ETA: 0s - loss: 0.0551 - accuracy: 0.9832WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 57s 30ms/step - loss: 0.0551 - accuracy: 0.9832

Epoch 6/10
1875/1875 [=====] - ETA: 0s - loss: 0.0462 - accuracy: 0.9857WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 58s 31ms/step - loss: 0.0462 - accuracy: 0.9857

Epoch 7/10
1874/1875 [=====>.] - ETA: 0s - loss: 0.0389 - accuracy: 0.9874WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 58s 31ms/step - loss: 0.0389 - accuracy: 0.9874

Epoch 8/10
1875/1875 [=====] - ETA: 0s - loss: 0.0337 - accuracy: 0.9895WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 57s 30ms/step - loss: 0.0337 - accuracy: 0.9895

Epoch 9/10
1875/1875 [=====] - ETA: 0s - loss: 0.0307 - accuracy: 0.9907WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 58s 31ms/step - loss: 0.0307 - accuracy: 0.9907

Epoch 10/10
1875/1875 [=====] - ETA: 0s - loss: 0.0271 - ac

```
curacy: 0.9916WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
1875/1875 [=====] - 57s 31ms/step - loss: 0.027
1 - accuracy: 0.9916
```

Out[27]: <tensorflow.python.keras.callbacks.History at 0x1950345d0d0>

Not much better than the baseline model but still good results, will an 99.+ percentage. with just the train data.

Testing on unseen data

```
In [28]: 1 test_labels_encoded = tf.one_hot(test_labels, 10) # One-hot encode test labels
2 test_loss, test_acc = model.evaluate(test_images, test_labels_encoded)
3 print('Test accuracy:', test_acc)

313/313 [=====] - 4s 12ms/step - loss: 0.0355 - accuracy: 0.9882
Test accuracy: 0.9882000088691711
```

Does well im glad overall with the final accuracy score of 0.98 Doesnt do better than training data, excellent loss score.

-

In []:

1